

CPSC 524 Assignment 4: CUDA Matrix Multiplication

Rami Pellumbi*

December 10, 2023

*M.S., Statistics & Data Science

1 Introduction

This assignment explores GPU programming using CUDA, centered around the task of matrix multiplication. The primary challenge involves constructing a CUDA kernel capable of multiplying two random rectangular matrices. This initial task serves as a gateway into the realms of parallel computing and efficient memory management, foundational elements in leveraging GPU architecture for computational tasks. As we progress, the report addresses more advanced techniques, such as the utilization of shared memory and the strategic optimization of thread computations, essential for enhancing the computational performance.

2 Project Organization

The project is laid out as follows:

```
├── docs/
├── bin/
├── include/
├── out/
├── src/
│   ├── serial.cu
│   └── task1.cu
├── build-run-task1a.sh
├── build-run-task1b.sh
├── build-run-task1c.sh
└── Makefile
```

- **docs/**: This folder contains LaTeX files and other documentation materials that pertain to the report.
- **bin/**: The **bin** folder holds compiled objects and executable files, centralizing the output of the compilation process.
- **include/**: Here, all the header files (**.h**) are stored.
- **out/**: The **out** folder stores the outputs from each task. It also houses the csv file containing data generated by the programs.
- **src/**: This directory houses the source files (**.cu**) that make up the benchmarks.
- **Shell Scripts**: The shell scripts are used to submit the job for the relevant task to slurm via **sbatch**.

3 Code Explanation, Compilation, and Execution

This section outlines the steps required to build and execute the code. The provided Bash scripts automate the entire process, making it straightforward to compile and run the code. All the below steps assume you are in the root of the project directory.

3.1 Automated Building and Execution

All related code is in the **src/** directory. There are multiple programs:

- **task1.cu**
- **task2.cu**
- **task3.cu**
- **task4.cu**

To run any one experiment, execute the relevant bash script, e.g., **build-run-task1a.sh**. It should be noted that for the **build-run-task1b.sh** the FP must be defined to double on line 1.

3.2 Post-Build Objects and Executables

Upon successful compilation and linking, an `obj/` subdirectory will be generated within the directory. This directory will contain the compiled output files. Additionally, the executable files for running each program will be situated in the `bin/` subdirectory.

3.3 Output Files From sbatch

The output files generated from running the code by submitting the relevant Bash script via `sbatch` will be stored in the `out` directory.

4 Task 1: CUDA Matrix Multiplication

First, a CUDA kernel is built to handle the multiplication of two random rectangular matrices: $C = AB$, where A is $n \times p$ and B is $p \times m$. Global memory is used to store the matrices and perform the computation. Each GPU thread computes only a single element of C . The input matrices A and B are stored row-wise into a one-dimensional array, and the output matrix C is stored in the same manner. In addition to the kernel, a host multiplication takes form using the *kij* access pattern:

```
for (int k = 0; k < p; k++) {
    size_t ia = i * p + k; // row i column k of A
    r = A[ia];
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < m; j++) {
            size_t ib = k * m + j; // row k column j of B
            size_t ic = i * m + j; // row i column j of C

            C[ic] += r * B[ib];
        }
    }
}
```

4.1 GPU Kernel Performance: float

The kernel was written to accept 7 arguments:

- `int n`: Number of rows in A and C .
- `int p`: Number of columns in A and rows in B .
- `int m`: Number of columns in B and C .
- `int block_dim_x`: The tile size to use for the computation.
- `int block_dim_y`: The tile size to use for the computation.
- `grid_dim_x`: The number of blocks to use in the x -direction.
- `grid_dim_y`: The number of blocks to use in the y -direction.

All device multiplication was validated in error performance against the serial code when able. This code was omitted from the individual task files for clarity, and the serial code put in its own `serial.cu` for clarity. The following bash script outline was used in Task 1, allowing for control over varying the block and grid size (handling the appropriate setting of grid dimensions):

```

# n,m,p blocks delimited by a space
sizes="1024,1024,1024 8192,8192,8192 1024,1024,8192 8192,8192,1024 8192,1024,8192"

for tuple in $sizes
do
    IFS=',' read -ra ADDR <<< "$tuple"
    n=${ADDR[0]}
    m=${ADDR[1]}
    p=${ADDR[2]}

    echo "Running n=$n, p=$p, m=$m"
    for blockx in 8 16 32 64 128 256
    do
        blocky=$((1024/$blockx))
        Grid_Dim_x=$((($m + $blockx - 1)/$blockx))
        Grid_Dim_y=$((($n + $blocky - 1)/$blocky))
        echo "Running block BLOCK_DIM_X=$blockx BLOCK_DIM_Y=$blocky"
        echo "With GRID_DIM_X=$Grid_Dim_x GRID_DIM_Y=$Grid_Dim_y"
        time ./bin/task1 $n $p $m $blockx $blocky $Grid_Dim_x $Grid_Dim_y
        time ./bin/task1 $n $p $m $blockx $blocky $Grid_Dim_x $Grid_Dim_y
        time ./bin/task1 $n $p $m $blockx $blocky $Grid_Dim_x $Grid_Dim_y
    done
done

```

Modifications of this script are carried over for parts 2 through 4. The performance table for the best performing (block_x, block_y) pair across the varying matrix sizes is:

Table 1: GPU Kernel Performance Float (ms)

(n, m, p)	(block x, blocky)	(grid x, grid y)	gpu time (ms)	cpu time (ms)
(1024,1024,1024)	(64, 16)	(16, 64)	3.628693	253.5439
(1024,1024,8192)	(128, 8)	(8, 128)	42.841162	2218.2016
(8192,8192,1024)	(16, 64)	(512, 128)	271.195058	28123.4727
(8192,1024,8192)	(64, 16)	(16, 512)	290.632660	25230.9069
(8192,8192,8192)	(16, 64)	(512, 128)	2179.101807	224239.7760

We observe that the GPU offers significant speedup, up to about 100× in some cases

4.2 GPU Kernel Performance: double

We assess the performance of the GPU on doubles for square matrix $n = m = p = 8192$. The results were:

Table 2: GPU Kernel Performance Double (ms)

(n,m,p)	(block x, block y)	(grid x, grid y)	gpu time
(8192,8192,8192)	(8,128)	(1024,64)	3462.878
(8192,8192,8192)	(16,64)	(512,128)	2870.899
(8192,8192,8192)	(32,32)	(256,256)	2826.502
(8192,8192,8192)	(64,16)	(128,512)	2944.203
(8192,8192,8192)	(128,8)	(64,1024)	3005.359
(8192,8192,8192)	(256,4)	(32,2048)	3231.829

We see that the best performing kernel was for block dimensions (32, 32) and grid dimensions (256, 256). This is a change from the best performing block dimensions of the float kernel, which was (16, 64). The gpu execution time increased from 2179.101807 to 2826.502, a 29.7% increase.

4.3 Maximum Matrix Size

The RTX 2080 Ti has 11,012 MegaBytes of memory. That is

$$11,012MB \times \frac{1024 KB}{MB} \times \frac{1024 B}{KB} = 11546918912 \text{ Bytes} := C.$$

Each float carries 4 bytes and thus the maximum matrix size that can be allocated is subject to the constraint

$$4(n \cdot p + p \cdot m + n \cdot m) \leq C.$$

For square matrices $n = p = m$, this reduces to

$$4 \cdot 3 \cdot n^2 \leq C \implies n \leq \sqrt{\frac{C}{12}} = 31020$$

Of course, this is the maximum size that can be allocated, but the actual maximum size is subject to the memory required by other processes. The maximum size that was able to be allocated was $n = 30536$. The timing result table is:

Table 3: Max Size GPU Kernel Performance Float (ms)

(n,m,p)	gpu time
(30536,30536,30536)	121486.5

5 Task 2: CUDA Matrix Multiplication with Shared Memory

Next, we modify the kernel to use shared memory. For simplicity, we use `block_dim_x = block_dim_y = tile_width`, but we do not assume that the matrix dimensions are multiples of the tile width. The algorithm details are implemented entirely based on Kirk and Hwu’s “Programming Massively Parallel Processors”. The best performing gpu time results are as follows (juxtaposed with the cpu time):

Table 4: Shared Memory GPU Kernel Performance Float (ms)

(n,m,p)	(block x, block y)	(grid x, grid y)	gpu time (ms)	cpu time (ms)
(1024,1024,1024)	(16,16)	(64,64)	2.154752	253.5439
(1024,1024,8192)	(16,16)	(64,64)	12.527520	2218.2016
(8192,8192,1024)	(16,16)	(512,512)	105.650368	28123.4727
(8192,1024,8192)	(16,16)	(64,512)	131.119908	25230.9069
(8192,8192,8192)	(32,32)	(256,256)	885.436859	224239.7760

We remark that there is a significant performance improvement from exploiting shared memory via tiled matrix multiplication. The best case tiled version significantly outperforms the best case non-tiled version.

6 Task 3: Reducing Tile Loads

Per the assignment, an important algorithmic decision in performance tuning is the granularity of thread computations. We attempt to improve performance by implementing a multi-tiled matrix multiplication kernel. The implementation was as follows:

```

__global__ void gpu_mult(FP *A, FP *B, FP *C, int n, int p, int m, int TILE_WIDTH)
{
    extern __shared__ FP tiles[];
    FP *Ads = &tiles[0]; // TILE_WIDTH x TILE_WIDTH
    FP *Bds[NTB];        // pointer to an array of NTB elements each of size TILE_WIDTH x TILE_WIDTH

    FP cvalues[NTB]; // initialize cvalues
    for (size_t kt = 0; kt < NTB; kt++)
    {
        // offset to appropriate shared memory location
        Bds[kt] = &tiles[(kt + 1) * TILE_WIDTH * TILE_WIDTH];
        cvalues[kt] = 0.;
    }

    int bx = blockIdx.x; int by = blockIdx.y;
    int tx = threadIdx.x; int ty = threadIdx.y;

    int row = by * TILE_WIDTH + ty; // col needs to be dynamically computed based on NTB idx
    int tile_idx = ty * TILE_WIDTH + tx;

    for (size_t ph = 0; ph < ceil((double)p / (double)TILE_WIDTH); ph++)
    {
        int col_bound_A = ph * TILE_WIDTH + tx;
        int row_bound_B = ph * TILE_WIDTH + ty;

        if (row < n && col_bound_A < p) // load Ads
        {
            int indexa = row * p + col_bound_A;
            Ads[tile_idx] = A[indexa];
        }

        for (size_t kt = 0; kt < NTB; kt++) // load multiple tiles of B into the Bds array
        {
            int col_offset = tx + NTB * blockDim.x * bx + kt * TILE_WIDTH;
            int indexb = row_bound_B * m + col_offset;

            if (col_offset < m && row_bound_B < p)
            {
                Bds[kt][tile_idx] = B[indexb];
            }
        }
        __syncthreads();

        for (size_t k = 0; k < TILE_WIDTH; k++)
            for (size_t kt = 0; kt < NTB; kt++)
            {
                cvalues[kt] += Ads[ty * TILE_WIDTH + k] * Bds[kt][k * TILE_WIDTH + tx];
            }
        __syncthreads();
    }

    for (size_t kt = 0; kt < NTB; kt++)
    {
        int col_offset = tx + NTB * blockDim.x * bx + kt * TILE_WIDTH;
        if (col_offset < m && row < n) C[row * m + col_offset] = cvalues[kt];
    }
}

```

The results for $n = m = p = 8192$ and a tile size of 32 across varying NTB values were:

Table 5: Multi Tiled GPU Kernel Performance Float (ms)

NTB	(n,m,p)	(block x, block y)	(grid x, grid y)	gpu time (ms)
1	(8192,8192,8192)	(32,32)	(256,256)	882.2599
2	(8192,8192,8192)	(32,32)	(256,256)	1123.4740
3	(8192,8192,8192)	(32,32)	(256,256)	1409.7855
4	(8192,8192,8192)	(32,32)	(256,256)	1659.7762
8	(8192,8192,8192)	(32,32)	(256,256)	2819.9041

I was unable to get performance to improve past NTB = 1 and thus concluded that the overhead of loading multiple tiles into shared memory was too great to overcome the performance gains of reducing the number of global memory loads.

7 Task 4: Handling All Tile Sizes

The task 3 program was implemented for the general case. Running it for $n = p = m = 1100$ and a tile size of 16×16 and 3 adjacent tiles yielded:

```
***Building task 4
rm -f obj/*.o bin/task1 bin/task4 bin/task2 bin/task3 bin/serial
nvcc -gencode=arch=compute_75,code=sm_75 -O3 -g -I include -o obj/serial.o -c src/serial.cu
Building binary bin/serial from object obj/serial.o
nvcc -gencode=arch=compute_75,code=sm_75 -lm -g -o bin/serial obj/serial.o
rm obj/serial.o
nvcc -gencode=arch=compute_75,code=sm_75 -O3 -g -I include -o obj/task4.o -c src/task4.cu
Building binary bin/task4 from object obj/task4.o
nvcc -gencode=arch=compute_75,code=sm_75 -lm -g -o bin/task4 obj/task4.o
rm obj/task4.o
Running n=1100, p=1100, m=1100
Running block BLOCK_DIM_X=16 BLOCK_DIM_Y=16
With GRID_DIM_X=69 GRID_DIM_Y=69
Device count = 1
Using device 0
Matrix Dimension = 1100
Block_Dim_x = 16, Block_Dim_y = 16, Grid_Dim_x = 69, Grid_Dim_y = 69
Time to calculate results on GPU: 4.602176 ms.
Time to calculate results on CPU: 315.723541 ms.
Approximate relative error between GPU and CPU: 4.315665e-08
```

We note that the relative error is very small, and the gpu time is significantly less than the cpu time.

8 Conclusion

In this report, we investigated the utilization of GPU programming using CUDA for matrix multiplication, revealing significant performance advantages over traditional CPU-based approaches. Our exploration highlighted the efficiency gains achievable through parallel processing inherent in GPU architecture. We further demonstrated how leveraging shared memory and tiling techniques can optimize GPU computations. Delving into advanced optimization with multi-tiled matrix multiplication, we observed that performance improvements are closely tied to specific GPU architectures. This finding emphasizes the necessity of understanding the nuances of GPU hardware for effective optimization.